

# Hadoop MapReduce

---

# What Is MapReduce?

---

- **MapReduce is a method for distributing a task across multiple nodes**
- **Each node processes data stored on that node**
  - Where possible
- **Consists of two phases:**
  - Map
  - Reduce

# MapReduce: Terminology

---


- A ***job*** is a ‘full program’
  - A complete execution of Mappers and Reducers over a dataset
  - In MapReduce 2, the term *application* is often used in place of ‘job’
- A ***task*** is the execution of a single Mapper or Reducer over a slice of data
- A ***task attempt*** is a particular instance of an attempt to execute a task
  - There will be at least as many task attempts as there are tasks
  - If a task attempt fails, another will be started by the JobTracker
  - *Speculative execution* (see later) can also result in more task attempts than completed tasks

# Hadoop Components: MapReduce

---


- **The Mapper**

- Each Map task (typically) operates on a single HDFS block
- Map tasks (usually) run on the node where the block is stored



- **Shuffle and Sort**

- Sorts and consolidates intermediate data from all mappers
- Happens as Map tasks complete and before Reduce tasks start



- **The Reducer**

- Operates on shuffled/sorted intermediate data (Map task output)
- Produces final output



## Example: Word Count

---

Input Data

```
the cat sat on the mat  
the aardvark sat on the sofa
```

Map


Reduce

Result

aardvark	1
cat	1
mat	1
on	2
sat	2
sofa	1
the	4

## Example: The WordCount Mapper

---




## Example: Shuffle & Sort

---

### Mapper Output


the	1
cat	1
sat	1
on	1
the	1
mat	1
the	1
aardvark	1
sat	1
on	1
the	1
sofa	1



### Intermediate Data






aardvark	1
cat	1
mat	1
on	1,1
sat	1,1
sofa	1
the	1,1,1,1

## Example: SumReducer



# Mappers Run in Parallel

- Hadoop runs Map tasks on the slave node where the block is stored (when possible)
  - Many Mappers can run in parallel
  - Minimizes network traffic



a	1,1,1,...
ac	1,1,1,1,1....
accumsan	1,1
ad	1,1,1,1....
adipiscing	1,1,1,1
aliquet	1,1,1,1....
...	

a	1,1,1,1....
ac	1,1,1,1,1...
ad	1,1,1,1,1...
aliquam	1,1,1
aliquet	1
auctor	1,1,1,1...
...	

a	1,1,1,...
ad	1,1,1,1....
aliquam	1,1,1
amet	1,1,1,1
blandit	1,1,1
...	

# MapReduce: The Mapper

---

- **Hadoop attempts to ensure that Mappers run on nodes which hold their portion of the data locally, to avoid network traffic**
  - Multiple Mappers run in parallel, each processing a portion of the input data
- **The Mapper reads data in the form of key/value pairs**
  - The Mapper may use or completely ignore the input key
  - For example, a standard pattern is to read one line of a file at a time
    - The key is the byte offset into the file at which the line starts
    - The value is the contents of the line itself
    - Typically the key is considered irrelevant
- **If the Mapper writes anything out, the output must be in the form of key/value pairs**

# MapReduce: The Reducer

---


- After the Map phase is over, all intermediate values for a given intermediate key are combined together into a list
- This list is given to a Reducer
  - There may be a single Reducer, or multiple Reducers
  - All values associated with a particular intermediate key are guaranteed to go to the same Reducer
  - The intermediate keys, and their value lists, are passed to the Reducer in sorted key order
  - This step is known as the ‘shuffle and sort’
- The Reducer outputs zero or more final key/value pairs
  - These are written to HDFS
  - In practice, the Reducer usually emits a single key/value pair for each input key

# Why Do We Care About Counting Words?

---

- **Word count is challenging over massive amounts of data**
  - Using a single compute node would be too time-consuming
  - Using distributed nodes requires moving data
  - Number of unique words can easily exceed available memory
    - Would need to store to disk
- **Statistics are simple aggregate functions**
  - Distributive in nature
  - e.g., max, min, sum, count
- **MapReduce breaks complex tasks down into smaller elements which can be executed in parallel**
- **Many common tasks are very similar to word count**
  - e.g., log file analysis

# Review: The MapReduce Flow



# A Sample MapReduce Program: WordCount

---

## Let's have a look at the code for WordCount

- This will demonstrate the Hadoop API

```
the cat sat on the mat  
the aardvark sat on the sofa
```



aardvark	1
cat	1
mat	1
on	2
sat	2
sofa	1
the	4

# Our MapReduce Program: WordCount

---

- To investigate the API, we will dissect the WordCount program
- This consists of three portions
  - The driver code
    - Code that runs on the client to configure and submit the job
  - The Mapper
  - The Reducer
- Before we look at the code, we need to cover some basic Hadoop API concepts

# Getting Data to the Mapper


---

- **The data passed to the Mapper is specified by an *InputFormat***
  - Specified in the driver code
  - Defines the location of the input data
    - Typically a file or directory
  - Determines how to split the input data into *input splits*
    - Each Mapper deals with a single input split
  - Creates a RecordReader object
    - RecordReader parses the input data into key/value pairs to pass to the Mapper

# Example: TextInputFormat

- **TextInputFormat**

- The default
- Creates LineRecordReader objects
- Treats each \n-terminated line of a file as a value
- Key is the byte offset of that line within the file



# Other Standard InputFormats

---

- **FileInputFormat**
  - Abstract base class used for all file-based InputFormats
- **KeyValueTextInputFormat**
  - Maps \n-terminated lines as ‘key [separator] value’
    - By default, [separator] is a tab
- **SequenceFileInputFormat**
  - Binary file of (key, value) pairs with some additional metadata
- **SequenceFileAsTextInputFormat**
  - Similar, but maps (key.toString(), value.toString())

## Keys and Values are Objects

---

- Keys and values in Hadoop are Java Objects
  - Not primitives
- Values are objects which implement `Writable`
- Keys are objects which implement `WritableComparable`

## What is Writable?

---

- The `Writable` interface makes serialization quick and easy for Hadoop
- Any value's type must implement the `Writable` interface
- Hadoop defines its own 'box classes' for strings, integers, and so on
  - `IntWritable` for ints
  - `LongWritable` for longs
  - `FloatWritable` for floats
  - `DoubleWritable` for doubles
  - `Text` for strings
  - Etc.

## What is WritableComparable?

---

- A **WritableComparable** is a **Writable** which is also **Comparable**
  - Two WritableComparables can be compared against each other to determine their ‘order’
  - Keys must be WritableComparables because they are passed to the Reducer in sorted order
  - We will talk more about WritableComparables later
- Note that despite their names, all Hadoop box classes implement both **Writable** and **WritableComparable**
  - For example, IntWritable is actually a WritableComparable

# The Driver: Complete Code

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.Job; import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class WordCount extends Configured implements Tool {

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new Configuration(), new WordCount(), args);
        System.exit(exitCode);
    }

    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            System.out.printf(
                "Usage: %s [generic options] <input dir> <output dir>\n",
                getClass().getSimpleName());
            return -1;
        }
        Job job = new Job(getConf());
        job.setJarByClass(WordCount.class);
        job.setJobName("Word Count");

        FileInputFormat.setInputPaths(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.setMapperClass(WordMapper.class);
        job.setReducerClass(SumReducer.class);
        job.setMapOutputKeyClass(Text.class);
        job.setMapOutputValueClass(IntWritable.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        boolean success = job.waitForCompletion(true);
        return success ? 0 : 1;
    }
}
```

# Driver Class Definition

```
public class WordCount extends Configured implements Tool {  
  
    public static void main(String[] args) {  
        int exitCode = ToolRunner.run(getConf(), WordCount.class, args);  
        System.exit(exitCode);  
    }  
  
    public int run(String[] args) throws Exception {  
  
        if (args.length != 2) {  
            System.out.printf(  
                "Usage: %s [generic options] <input dir> <output dir>\n", getClass().getSimpleName());  
            return -1;  
        }  
        Job job = new Job(getConf());  
        job.setJarByClass(WordCount.class); job.setJobName("Word Count");  
  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.setMapperClass(WordMapper.class);  
        job.setReducerClass(SumReducer.class);  
        job.setMapOutputKeyClass(Text.class);  
        job.setMapOutputValueClass(IntWritable.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        boolean success = job.waitForCompletion(true);  
        return success ? 0 : 1;  
    }  
}
```

The driver class implements the Tool interface and extends the Configured class.

# Main Method

```
public class WordCount extends Configured implements Tool {  
  
    public static void main(String[] args) throws Exception {  
        int exitCode = ToolRunner.run(new Configuration(),  
            new WordCount(), args);  
        System.exit(exitCode);  
    }  
  
    public int run(String[] args) throws IOException, InterruptedException,  
        ClassNotFoundException {  
        if (args.length != 2)  
            System.out.printf("Usage: %s [input path] [output path]\n",  
                getClass().getSimpleName());  
        return -1;  
    }  
    Job job = new Job(getConf());  
    job.setJarByClass(WordCount.class);  
    job.setJobName("Word Count");  
  
    FileInputFormat.setInputPaths(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    job.setMapperClass(WordMapper.class);  
    job.setReducerClass(SumReducer.class);  
    job.setMapOutputKeyClass(Text.class);  
    job.setMapOutputValueClass(IntWritable.class);  
  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
  
    ...  
}
```

The driver main method calls ToolRunner.run.

## Run Method

```
public class WordCount {  
    public static void main(String[] args) {  
        int exitCode = ToolRunner.run(new WordCount(), args);  
        System.exit(exitCode);  
    }  
}
```

The driver `run` method creates, configures, and submits the job.

```
public int run(String[] args) throws Exception {  
    if (args.length != 2) {  
        System.out.printf(  
            "Usage: %s [generic options] <input dir> <output dir>\n", getClass().getSimpleName());  
        return -1;  
    }  
    Job job = new Job(getConf());  
    job.setJarByClass(WordCount.class);  
    job.setJobName("Word Count");  
  
    FileInputFormat.setInputPaths(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    ...  
}
```

## Creating a New Job Object

---

- **The Job class allows you to set configuration options for your MapReduce job**
  - The classes to be used for your Mapper and Reducer
  - The input and output directories
  - Many other options
- **Any options not explicitly set in your driver code will be read from your Hadoop configuration files**
  - Usually located in /etc/hadoop/conf
- **Any options not specified in your configuration files will use Hadoop's default values**
- **You can also use the Job object to submit the job, control its execution, and query its state**

## Configuring the Job: Specifying the InputFormat

---

- The default InputFormat (`TextInputFormat`) will be used unless you specify otherwise
- To use an InputFormat other than the default, use e.g.

```
job.setInputFormatClass (KeyValueTextInputFormat.class)
```

## Configuring the Job: Determining Which Files To Read

---

- By default, `FileInputFormat.setInputPaths()` will read all files from a specified directory and send them to Mappers
  - Exceptions: items whose names begin with a period (.) or underscore (\_)
  - Globs can be specified to restrict input
    - For example, `/2010/*/01/*`
- Alternatively, `FileInputFormat.addInputPath()` can be called multiple times, specifying a single file or directory each time
- More advanced filtering can be performed by implementing a `PathFilter`
  - Interface with a method named `accept`
    - Takes a path to a file, returns `true` or `false` depending on whether or not the file should be processed

## Configuring the Job: Specifying Final Output With OutputFormat

---

- **FileOutputFormat.setOutputPath()** specifies the directory to which the Reducers will write their final output
- The driver can also specify the format of the output data
  - Default is a plain text file
  - Could be explicitly written as  
`job.setOutputFormatClass(TextOutputFormat.class)`
- We will discuss OutputFormats in more depth in a later chapter

# Configuring the Job: Specifying the Mapper and Reducer Classes

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
        Job job = new Job();  
        job.setJarByClass(WordCount.class);  
        job.setJobName("Word Count");  
  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.setMapperClass(WordMapper.class);  
        job.setReducerClass(SumReducer.class);  
    }  
}
```

Give the Job object information about which classes are to be instantiated as the Mapper and Reducer.


```
    boolean success = job.waitForCompletion(true);  
    System.exit(success ? 0 : 1);  
}
```

# Default Mapper and Reducer Classes


---

- Setting the Mapper and Reducer classes is optional
- If not set in your driver code, Hadoop uses its defaults

- IdentityMapper



- IdentityReducer



# Configuring the Job: Specifying the Intermediate Data Types

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
        Job job = new Job();  
        job.setJarByClass(WordCount.class);  
        job.setJobName("Word Count");  
  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.setMapperClass(WordMapper.class);  
        job.setReducerClass(SumReducer.class);  
  
        job.setMapOutputKeyClass(Text.class);  
        job.setMapOutputValueClass(IntWritable.class);  
    }  
}
```

Specify the types for the intermediate output keys and values produced by the Mapper.

# Configuring the Job: Specifying the Final Output Data Types

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
        Job job = new Job();  
        job.setJarByClass(WordCount.class);  
        job.setJobName("Word Count");  
  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.setMapperClass(WordMapper.class);  
        j  
        j  
        Specify the types for the Reducer's output keys and values.  
        j  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        boolean success = job.waitForCompletion(true);  
        System.exit(success ? 0 : 1);  
    }  
}
```

# Running The Job (1)

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.out.printf("Usage: WordCount <input dir> <output dir>\n");  
            System.exit(-1);  
        }  
        Job job = new Job();  
        job.setJarByClass(WordCount.class);  
        job.setJobName("Word Count");  
  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
    }  
}
```

j  
j Start the job and wait for it to complete. Parameter is a Boolean,  
j specifying verbosity: if true, display progress to the user. Finally, exit  
j with a return code.  
j

```
boolean success = job.waitForCompletion(true);  
System.exit(success ? 0 : 1);  
}
```

## Running The Job

---

- **There are two ways to run your MapReduce job:**
  - `job.waitForCompletion()`
    - Blocks (waits for the job to complete before continuing)
  - `job.submit()`
    - Does not block (driver code continues as the job is running)
- **The client determines the proper division of input data into InputSplits, and then sends the job information to the JobTracker daemon on the cluster**

# ToolRunner Command Line Options

---

- ToolRunner allows the user to specify configuration options on the command line
- Commonly used to specify Hadoop properties using the **-D** flag
  - Will override any default or site properties in the configuration
  - But will *not* override those set in the driver code

```
$ hadoop jar myjar.jar MyDriver \
-D mapred.reduce.tasks=10 myinputdir myoutputdir
```

- Note that **-D** options must appear before any additional program arguments
- Can specify an XML configuration file with **-conf**
- Can specify the default filesystem with **-fs uri**
  - Shortcut for **-D fs.default.name=uri**

# The Mapper: Complete Code

---

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class WordMapper extends Mapper<LongWritable, Text, Text,
IntWritable> {
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        String line = value.toString();

        for (String word : line.split("\W+")) {
            if (word.length() > 0) {

                context.write(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

# The Mapper: Class Declaration (1)

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
        String line = value.toString();

        for (String word : line.split("\\w+")) {
            if (word.length() > 0) {

                context.write(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

Mapper classes extend the Mapper base class.

## The Mapper: Class Declaration (2)

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String word = value.toString();
        for (int i = 0; i < word.length(); i++) {
            if (Character.isLetter(word.charAt(i))) {
                context.write(new Text(word.substring(0, i + 1)), new IntWritable(1));
            }
        }
    }
}
```

Specify generic types that declare four type parameters: the input key and value types, and the output (intermediate) key and value types. Keys must be WritableComparable, and values must be Writable.

Input key and  
value types

Intermediate key  
and value types

# The map Method

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String word = value.toString();
        for (String w : words) {
            if (w.equals(word)) {
                context.write(key, value);
            }
        }
    }
}
```

The map method is passed a key, a value, and a Context object. The Context is used to write the intermediate data. It also contains information about the job's configuration.

# The map Method: Processing The Line (1)

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = value.toString();

        for (int i = 0; i < line.length(); i++) {
            char c = line.charAt(i);
            if (Character.isLetter(c)) {
                word.set(c);
                context.write(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

value is a Text object, so retrieve the string it contains.

## The map Method: Processing The Line (2)

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = value.toString();
        for (String word : line.split("\\w+")) {
            if (word.length() > 0) {
                cont
            }
        }
    }
}
```

Split the string up into words using a regular expression with non-alphanumeric characters as the delimiter, and then loop through the words.

# The map Method: Outputting Intermediate Data

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
```

```
@Override
public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
```

To emit a (key, value) pair, call the `write` method of the Context object. The key will be the word itself, the value will be the number 1. Recall that the output key must be a `WritableComparable`, and the value must be a `Writable`.

```
    context.write(new Text(word), new IntWritable(1));
}
```

```

}
}
```

# The Reducer: Complete Code

---

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable>
{
    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int wordCount = 0;
        for (IntWritable value : values) {
            wordCount += value.get();
        }
        context.write(key, new IntWritable(wordCount));
    }
}
```

# The Reducer: Class Declaration

```
public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable>
{
    @Override
    public void reduce(Text key, Iterable<IntWritable> values,
                      OutputCollector<Text, IntWritable> output)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        output.collect(key, new IntWritable(sum));
    }
}
```

Reducer classes extend the `Reducer` base class. The four generic type parameters are: input (intermediate) key and value types, and final output key and value types.

Intermediate key and value types

Output key and value types

# The reduce Method

```
public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable>
{
    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int wordCount = 0;
        for (IntWritable value : values) {
            wordCount += value.get();
        }
        context.write(key, new IntWritable(wordCount));
    }
}
```

The `reduce` method receives a key and an `Iterable` collection of objects (which are the values emitted from the Mappers for that key); it also receives a `Context` object.

## The reduce Method: Processing The Values

```
public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable>
{
    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int wordCount = 0;
        for (IntWritable value : values) {
            wordCount += value.get();
        }
    }
}
```

We use the Java for-each syntax to step through all the elements in the collection. In our example, we are merely adding all the values together. We use `value.get()` to retrieve the actual numeric value each time.

# The reduce Method: Writing The Final Output

```
public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable>
{
    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int wordCount = 0;
        for (IntWritable value : values) {
            wordCount += value.get();
        }
        context.write(key, new IntWritable(wordCount));
    }
}
```

Finally, we write the output key-value pair to HDFS using the `write` method of our `Context` object.

## Ensure Types Match (1)

- Mappers and Reducers declare input and output type parameters
- These must match the types used in the class

Input key and value types

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        ...
        context.write(new Text(word), new IntWritable(1));
        ...
    }
}
```

Output key and value types

## Ensure Types Match (2)

- Output types must also match those set in the driver

```
public class WordMapper extends Mapper<LongWritable,  
Text, Text, IntWritable> {  
    ...  
}
```

Mapper

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        ...  
        job.setMapOutputKeyClass(Text.class);  
        job.setMapOutputValueClass(IntWritable.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        ...  
    }  
}
```

driver code

```
public class SumReducer extends Reducer<Text,  
IntWritable, Text, IntWritable> {  
    ...  
}
```

Reducer

## What Is The Old API?

---

- When Hadoop 0.20 was released, a ‘New API’ was introduced
  - Designed to make the API easier to evolve in the future
  - Favors abstract classes over interfaces
- Some developers still use the Old API
  - Until CDH4, the New API was not absolutely feature-complete
- All the code examples in this course use the New API

# New API vs. Old API: Some Key Differences (1)

New API	Old API
<pre>import org.apache.hadoop.mapreduce.*</pre>	<pre>import org.apache.hadoop.mapred.*</pre>
<b>Driver code:</b>  Configuration conf = new Configuration(); Job job = new Job(conf); job.setJarByClass(Driver.class); job.setSomeProperty(...); ... job.waitForCompletion(true);	<b>Driver code:</b>  JobConf conf = new JobConf(Driver.class); conf.setSomeProperty(...); ... JobClient.runJob(conf);
<b>Mapper:</b>  public class MyMapper extends Mapper {  public void map(Keytype k, Valuetype v, Context c) { ... c.write(key, val); } }	<b>Mapper:</b>  public class MyMapper extends MapReduceBase implements Mapper {  public void map(Keytype k, Valuetype v, OutputCollector o, Reporter r) { ... o.collect(key, val); } }

## New API vs. Old API: Some Key Differences (2)

New API	Old API
<b>Reducer:</b> <pre>public class MyReducer extends Reducer {      public void reduce(Keytype k,                       Iterable&lt;Valuetype&gt; v, Context c) {         for(Valuetype eachval : v) {             // process eachval             c.write(key, val);         }     } }</pre>	<b>Reducer:</b> <pre>public class MyReducer extends MapReduceBase     implements Reducer {      public void reduce(Keytype k,                       Iterator&lt;Valuetype&gt; v,                       OutputCollector o, Reporter r) {         while(v.hasNext()) {             // process v.next()             o.collect(key, val);         }     } }</pre>
setup(Context c) <b>(See later)</b>	configure(JobConf job)
cleanup(Context c) <b>(See later)</b>	close()

## MRv1 vs MRv2, Old API vs New API

---

- There is a lot of confusion about the New and Old APIs, and MapReduce version 1 and MapReduce version 2
- The chart below should clarify what is available with each version of MapReduce

	Old API	New API
MapReduce v1	✓	✓
MapReduce v2	✓	✓

- Summary: Code using either the Old API or the New API will run under MRv1 and MRv2

## The setup Method

---

- It is common to want your Mapper or Reducer to execute some code before the `map` or `reduce` method is called for the first time
  - Initialize data structures
  - Read data from an external file
  - Set parameters
- The `setup` method is run before the `map` or `reduce` method is called for the first time

```
public void setup(Context context)
```

## The cleanup Method

---

- Similarly, you may wish to perform some action(s) after all the records have been processed by your Mapper or Reducer
- The `cleanup` method is called before the Mapper or Reducer terminates

```
public void cleanup(Context context) throws  
    IOException, InterruptedException
```

## Passing Parameters

---

```
public class MyDriverClass {  
    public int main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        conf.setInt ("paramname", value);  
        Job job = new Job(conf);  
        ...  
        boolean success = job.waitForCompletion(true);  
        return success ? 0 : 1;  
    }  
}
```

```
public class MyMapper extends Mapper {  
  
    public void setup(Context context) {  
        Configuration conf = context.getConfiguration();  
        int myParam = conf.getInt("paramname", 0);  
        ...  
    }  
    public void map...  
}
```

# Reuse of Objects is Good Practice (1)

- It is generally good practice to reuse objects
  - Instead of creating many new objects
- Example: Our original WordCount Mapper code

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String word = value.toString();
        for (String line : lines) {
            if (line.contains(word)) {
                context.write(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

Each time the `map()` method is called, we create a new `Text` object and a new `IntWritable` object.

## Reuse of Objects is Good Practice (2)

- Instead, this is better practice:

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text wordObject = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
        String line = value.toString();

        for (String word : line.split("\\\\w+")) {
            if (word.length() > 0) {
                wordObject.set(word);
                context.write(wordObject, one);
            }
        }
    }
}
```

Create objects for the key and value outside of your `map()` method

## Reuse of Objects is Good Practice (3)

- Instead, this is better practice:

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text wordObject = new Text();
```

```
@Override
public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
```

Within the `map()` method, populate the objects and write them out. Hadoop will take care of serializing the data so it is perfectly safe to re-use the objects.

```
    if (word.length() > 0) {
        wordObject.set(word);
        context.write(wordObject, one);
    }
}
}
```

## Object Reuse: Caution!

---

- Hadoop re-uses objects all the time
- For example, each time the Reducer is passed a new value, the same object is reused
- This can cause subtle bugs in your code
  - For example, if you build a list of value objects in the Reducer, each element of the list will point to the same underlying object
    - Unless you do a deep copy

## Map-Only MapReduce Jobs

---

- **There are many types of job where only a Mapper is needed**
- **Examples:**
  - Image processing
  - File format conversion
  - Input data sampling
  - ETL

## Creating Map-Only Jobs

---

- To create a Map-only job, set the number of Reducers to 0 in your Driver code

```
job.setNumReduceTasks(0);
```

- Call the Job.setOutputKeyClass and Job.setOutputValueClass methods to specify the output types
  - Not the Job.setMapOutputKeyClass and Job.setMapOutputValueClass methods
- Anything written using the Context.write method in the Mapper will be written to HDFS
  - Rather than written as intermediate data
  - One file per Mapper will be written

## How Many Reducers Do You Need?

---

- An important consideration when creating your job is to determine the number of Reducers specified
- Default is a single Reducer
- With a single Reducer, one task receives *all* keys in sorted order
  - This is sometimes advantageous if the output must be in completely sorted order
  - Can cause significant problems if there is a large amount of intermediate data
    - Node on which the Reducer is running may not have enough disk space to hold all intermediate data
    - The Reducer will take a long time to run

## Jobs Which Require a Single Reducer

---

- If a job needs to output a file where all keys are listed in sorted order, a single Reducer must be used
- Alternatively, the **TotalOrderPartitioner** can be used
  - Uses an externally generated file which contains information about intermediate key distribution
  - Partitions data such that all keys which go to the first Reducer are smaller than any which go to the second, etc
  - In this way, multiple Reducers can be used
  - Concatenating the Reducers' output files results in a totally ordered list

## Jobs Which Require a Fixed Number of Reducers

---

- Some jobs will require a specific number of Reducers
- Example: a job must output one file per day of the week
  - Key will be the weekday
  - Seven Reducers will be specified
  - A Partitioner will be written which sends one key to each Reducer

## Jobs With a Variable Number of Reducers (1)

---

- **Many jobs can be run with a variable number of Reducers**
- **Developer must decide how many to specify**
  - Each Reducer should get a reasonable amount of intermediate data, but not too much
  - Chicken-and-egg problem
- **Typical way to determine how many Reducers to specify:**
  - Test the job with a relatively small test data set
  - Extrapolate to calculate the amount of intermediate data expected from the ‘real’ input data
  - Use that to calculate the number of Reducers which should be specified

## Jobs With a Variable Number of Reducers (2)

---

- **Note: you should take into account the number of Reduce slots likely to be available on the cluster**
  - If your job requires one more Reduce slot than there are available, a second ‘wave’ of Reducers will run
    - Consisting just of that single Reducer
    - Potentially doubling the amount of time spent on the Reduce phase
  - In this case, increasing the number of Reducers further may cut down the time spent in the Reduce phase
    - Two or more waves will run, but the Reducers in each wave will have to process less data